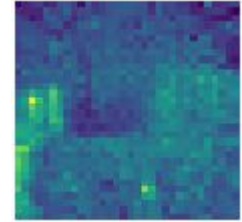
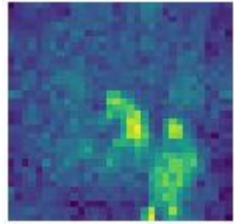


# IAS0360 Final Project - Topic 2:

# Thermal Image Classification



**presented by: Felix, Omar, Tobias**

# Agenda

- **General Workflow**
- **Different Model Approaches**
  - **Semi-supervised clustering**
  - **Squeezenet**
  - **Gradient Boosted Trees**
- **Conclusion**

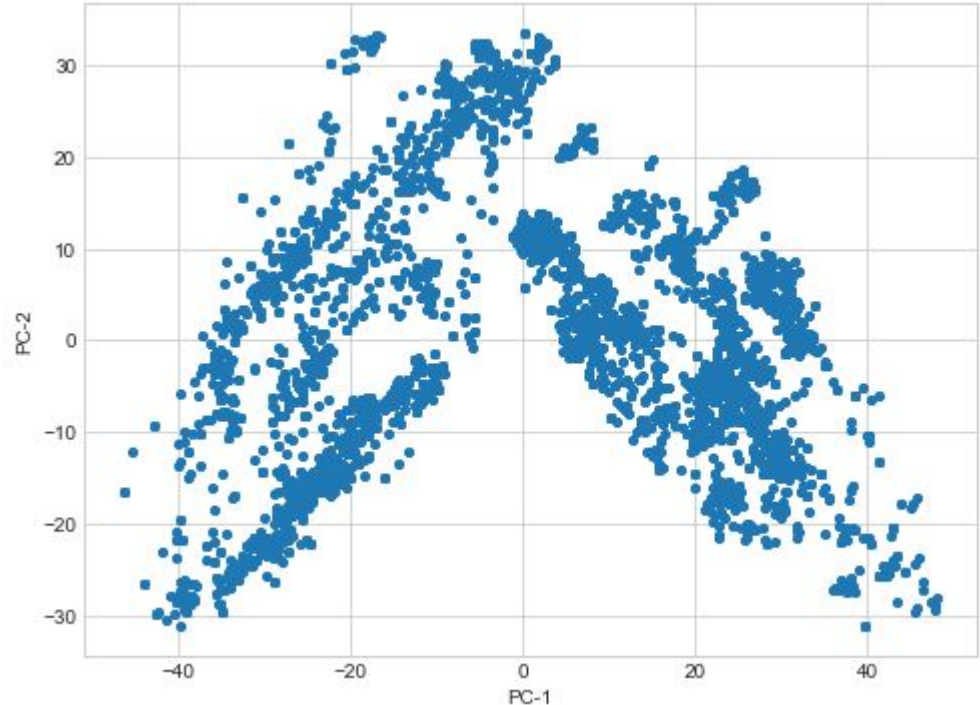
# Workflow - CRISP DM

Cross Industry Standard Process for Data Mining

- Business Understanding
- Data Understanding
- Data Preparation
- Modelling
- Evaluation
- Deployment


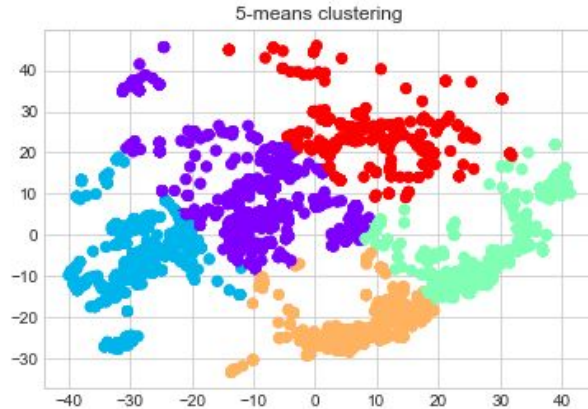
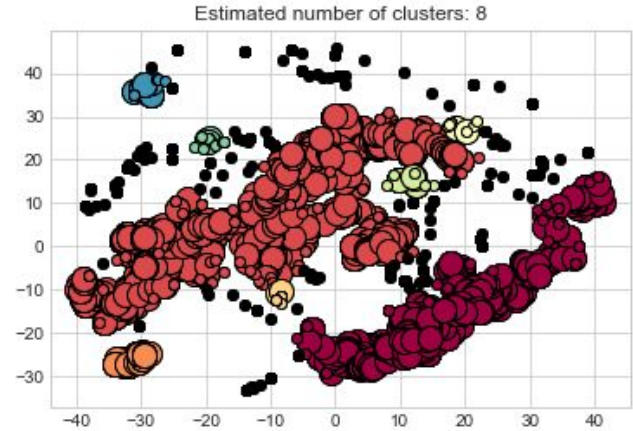
# Data Exploration (1)

- C088 & 3078 sensor
  - 32x32 dimensions
  - PCA-2
  - 35% explained



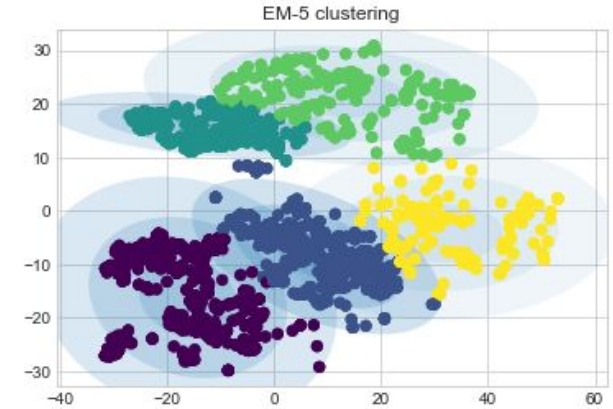

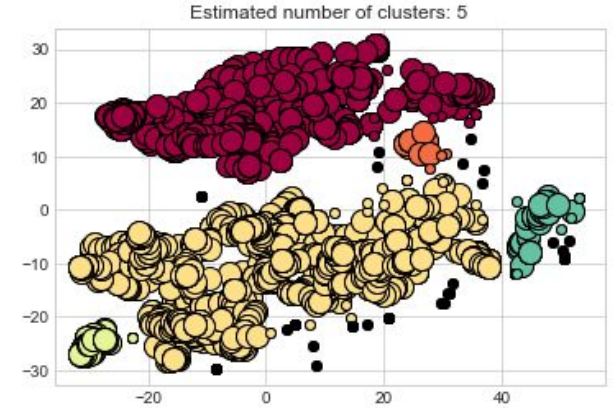
# Data Exploration (2)

- Exploring sensor C088:
  - DBSCAN
  - EM
  - K-means



# Data Exploration (3)

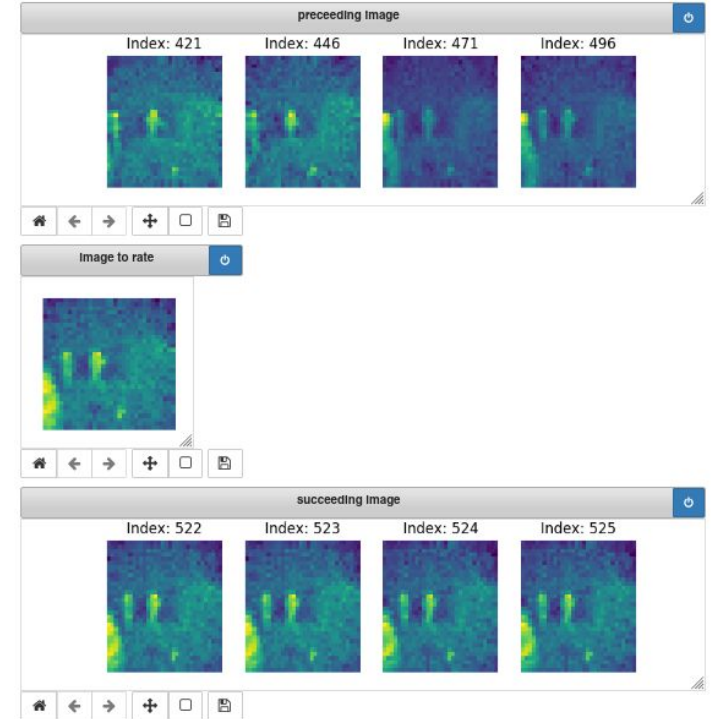
- Exploring sensor 3078:
  - DBSCAN
  - EM
  - K-means



# Labelling

Human	Several humans	Non-human heat	Ambiguous heat	No heat
prev	next			

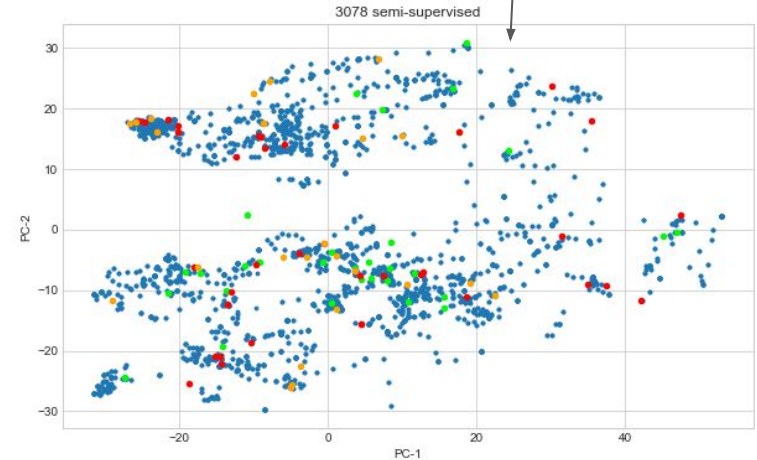
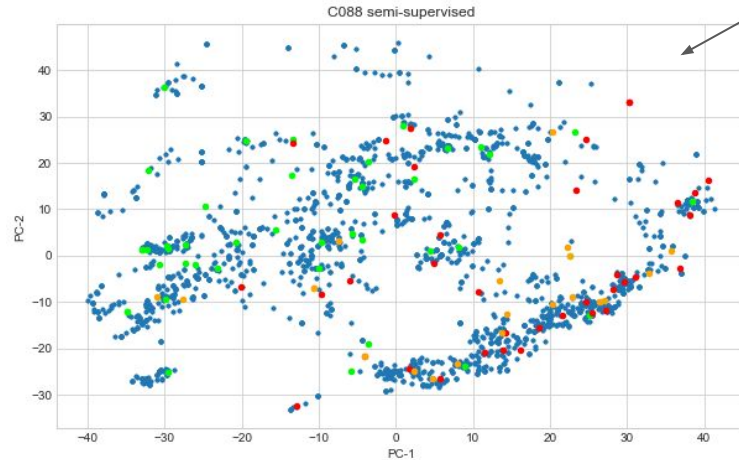
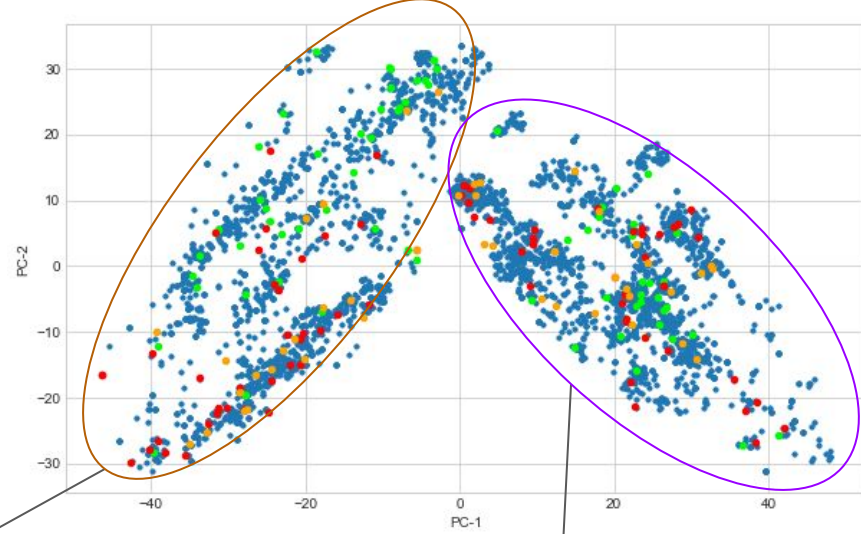
- Jupyter Notebook  
+ Pidgeon XT  
<https://bit.ly/3qOtV2n>
- Splits:  
4 x 50 = 200 samples
- Labelled
  - 1 individually each
  - 1 badge collaboratively





# Data Exploration (4)

- Semi-supervised
  - Human | Several humans (green)
  - Non-human | No heat (red)
  - Ambiguous (orange)
- **DBSCAN** yields best insights



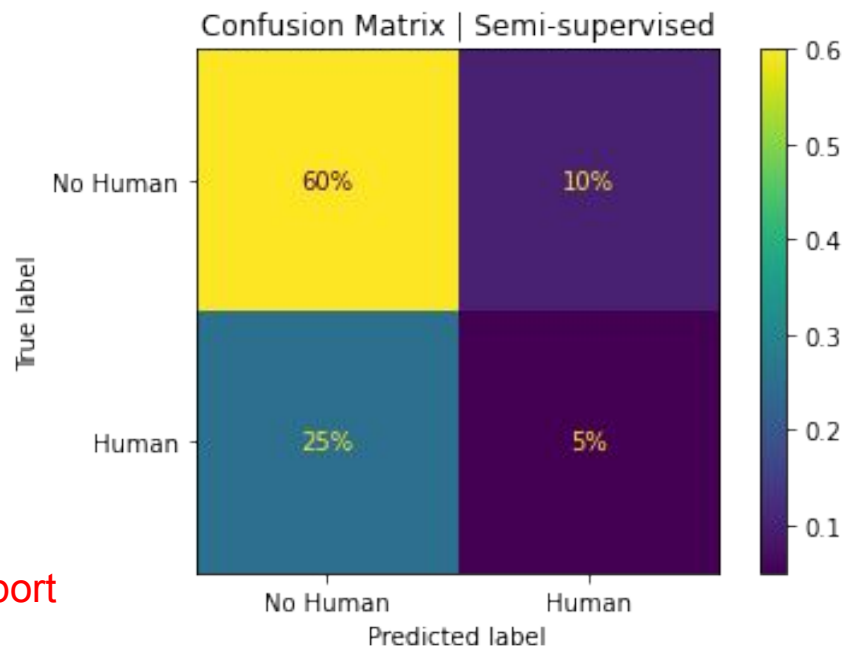


# Semi-Supervised Classification

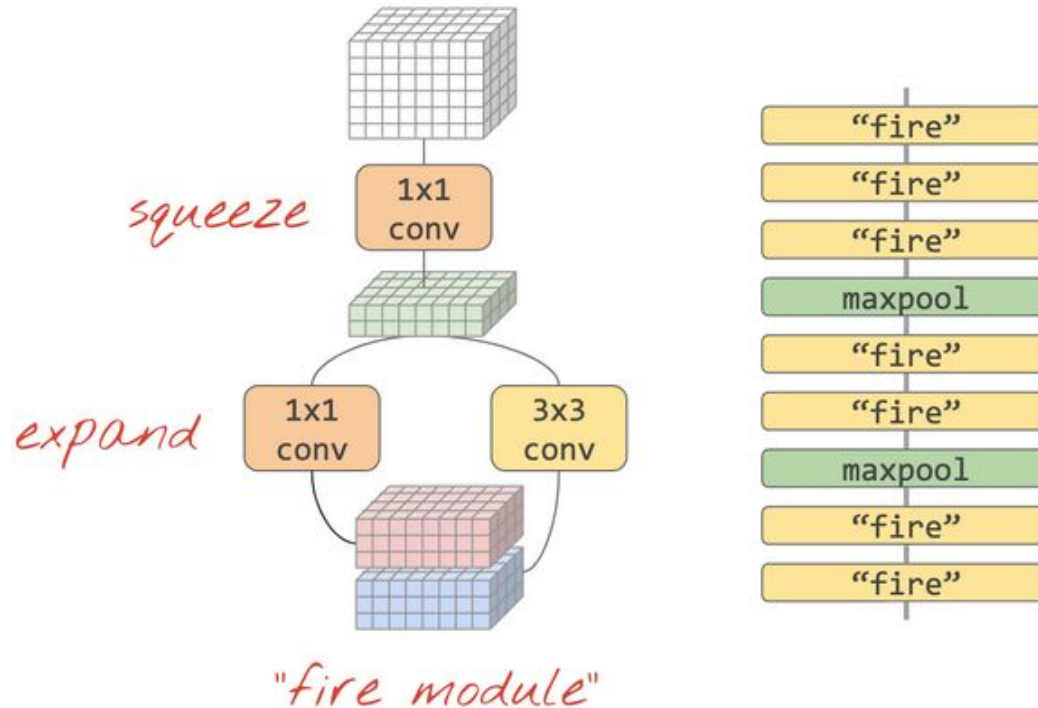
- Semi-supervised - Test Accuracy: 65 %
  - 180 Training, 20 Test Samples
- Self-trained classifier
  - Support Vector Classifier (SVC)
  - Other models possible (OUT OF SCOPE)

**Train 180 → Convert 10.000 → Predict 20**

**Problems: Bias in 180 labels & No sklearn/ONNX support**



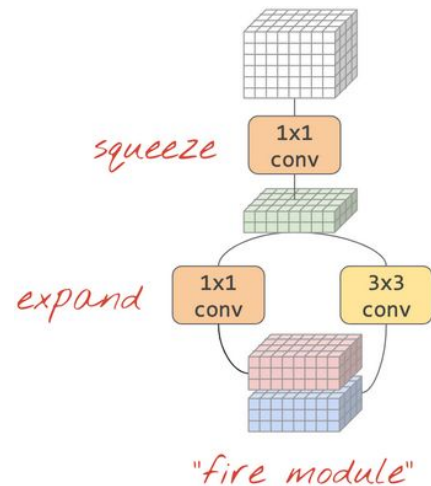
# Squeezenet - CNN Approach



<https://bit.ly/3nB1lze>

# Squeezenet - Fire Module

- Leveraging Keras  
**functional** model approach:



```
def fire(x, squeeze, expand):  
    y = tf.keras.layers.Conv2D(filters=squeeze, kernel_size=1, activation='relu', padding='same')(x)  
    y = tf.keras.layers.BatchNormalization(momentum=bnmomentum)(y)  
    y1 = tf.keras.layers.Conv2D(filters=expand//2, kernel_size=1, activation='relu', padding='same')(y)  
    y1 = tf.keras.layers.BatchNormalization(momentum=bnmomentum)(y1)  
    y3 = tf.keras.layers.Conv2D(filters=expand//2, kernel_size=3, activation='relu', padding='same')(y)  
    y3 = tf.keras.layers.BatchNormalization(momentum=bnmomentum)(y3)  
    return tf.keras.layers.concatenate([y1, y3])
```

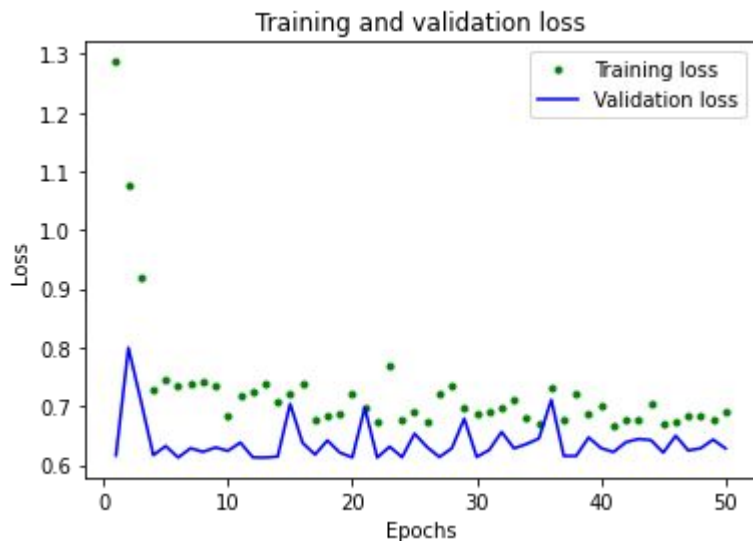
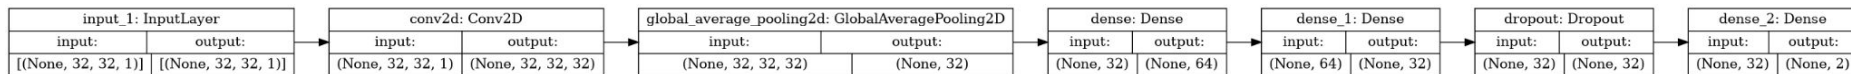
<https://bit.ly/3nB1lze>

# Squeezenet - Complete Model

```
x = tf.keras.layers.Input(shape=[*IMAGE_SIZE,1]) # input is 32x32 pixels RGB
y = tf.keras.layers.Conv2D(kernel_size=3, filters=32, padding='same', use_bias=True, activation='relu')(x)
y = tf.keras.layers.BatchNormalization(momentum=bnmomentum)(y)
y = tf.keras.layers.Dropout(0.2)(y)
y = fire_module(24, 48)(y)
y = tf.keras.layers.MaxPooling2D(pool_size=2)(y)
y = fire_module(48, 96)(y)
y = tf.keras.layers.MaxPooling2D(pool_size=2)(y)
y = fire_module(64, 128)(y)
y = tf.keras.layers.MaxPooling2D(pool_size=2)(y)
y = fire_module(48, 96)(y)
y = tf.keras.layers.MaxPooling2D(pool_size=2)(y)
y = fire_module(24, 48)(y)
y = tf.keras.layers.GlobalAveragePooling2D()(y)
y = tf.keras.layers.Dense(2, activation='softmax')(y)
```


<https://bit.ly/3nB1lze>

# Baseline CNN



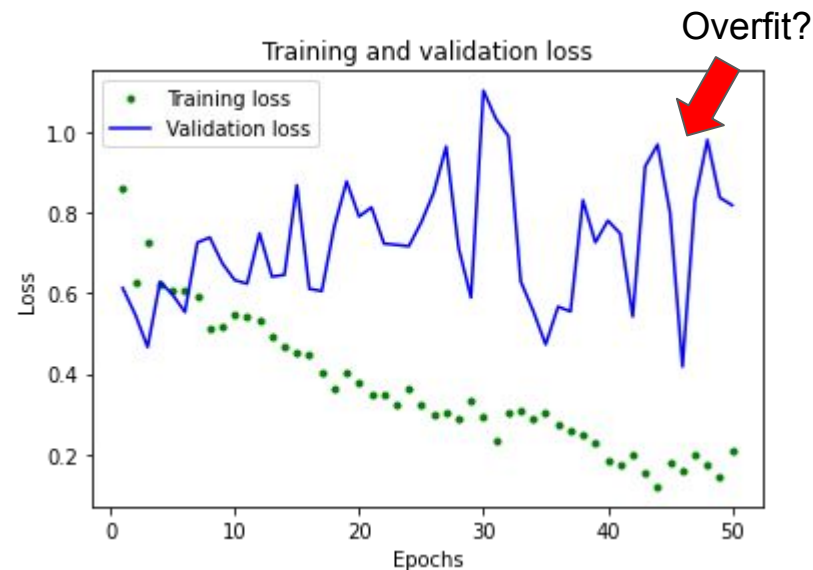

Epoch 50/50 - 0s 29ms/step  
 loss: 0.6900 - **accuracy: 0.6062**  
 val\_loss: 0.6282 - val\_accuracy: 0.7000

Test loss: 0.62011 / **Test accuracy: 0.69999**



Lowpass Threshold at 15.0

# Squeezenet CNN



Epoch 50/50 1s 51ms/step  
loss: 0.1196 - **accuracy: 0.9563**  
val\_loss: 0.6548 - val\_accuracy: 0.8000

Test loss: 0.79038 / **Test accuracy: 0.75**

# Squeezenet Implementation (1)

- Issues with STM32 'shrinking tool'

```
# stm32ai quantize -q cfg_squeezenet-binary.json
```

Neural Network Tools for STM32AI v1.5.1 (STM.ai v7.0.0-RC8)

```
-- Loading/Initializing dataset
```

```
X_test (20, 32, 32, 1)
```

```
y_test (20, 2)
```

```
-- Loading/Initializing dataset - done (elapsed time 0.002s)
```

```
-- Loading original model
```

```
-- Loading original model - done (elapsed time 0.457s)
```

```
-- Testing original model
```

```
Original model - test accuracy (loss): 0.75 (0.202)
```

```
-- Testing original model - done (elapsed time 0.664s)
```

**The original network has several branches.**

**We don't support quantization of these topologies.**

Process finished.

```
Total params: 119,890
Trainable params: 118,578
Non-trainable params: 1,312
```

1.7M squeezenet-binary.h5




Too big for MCU



# Squeezenet Implementation (2)

- TF lite (micro) with quantization



The screenshot shows the STM32Cube.AI web interface. At the top, there are tabs for 'Main', 'network', and a '+' icon. Under the 'network' tab, the 'Model inputs' section shows 'network' as the input. Below this, a dropdown menu is set to 'TFLite', and another dropdown menu is set to 'TFLite Micro runtime', which is highlighted by a red arrow. Below these, there are two 'Browse...' buttons for selecting a model file. The 'Model' field shows the path 'nal-project/cnn/models/squeezenet-binary.tflite'. A black bar with the text 'Please wait...' is visible below the model selection area.

Analyzing Network

Analyzing model  
/home/weiss/STM32Cube/Repository/Packs/STMicroelectronics/X-CUBE-AI/7.0.0  
/Utilities/linux/stm32tfml /home/weiss/git/ias0360-final-project/cnn/models/squeezenet-binary..  
Flash:160696

**Analyze complete on AI model**

Ram:114944

Model	Size
TensorFlow	2996689 bytes
TensorFlow Lite	494608 bytes (reduced by 2502081 bytes)
TensorFlow Lite Quantized	160696 bytes (reduced by 333912 bytes)

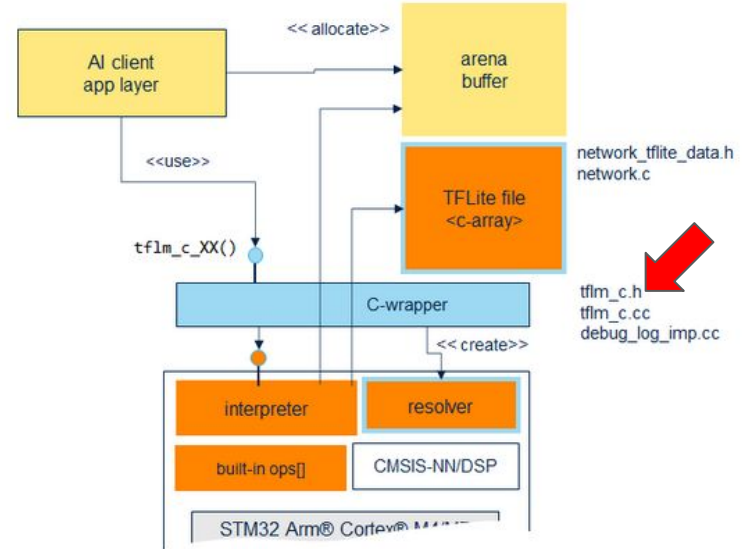
# Squeezenet Implementation (3)

- Local CubeMX AI 7.0.0

## Documentation:

~/STMicroelectronics/X-CUBE-AI/7.0.0/Documentation/tflite\_micro\_support.html#x-cube-ai-solution

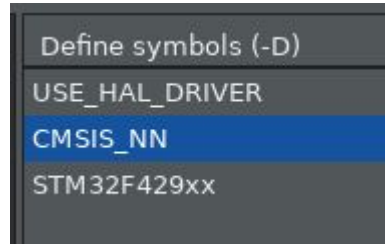
```
Include paths (-I)
./Inc
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party/ruy"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party/cmsis/CMSIS/DSP/Include"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party/cmsis"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party/gemmlowp"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party/cmsis/CMSIS/NN/Include"
"${workspace_loc}/${ProjName}/Middlewares/tensorflow/third_party/flatbuffers/include"
"${workspace_loc}/${ProjName}/Drivers/CMSIS/Include"
"${workspace_loc}/${ProjName}/Drivers/BSP/STM32F429I-Discovery"
"${workspace_loc}/${ProjName}/Drivers/BSP/Components/Common"
./Drivers/STM32F4xx_HAL_Driver/Inc
./Drivers/STM32F4xx_HAL_Driver/Inc/Legacy
${workspace_loc}/${ProjName}/Drivers/CMSIS/Device/ST/STM32F4xx/Include}
```



# Squeezenet Implementation (4)

- CMSYS Fast MCU NN Kernels:

“New CMSIS-NN Neural Network Kernels  
Boost Efficiency in Microcontrollers by ~5x”

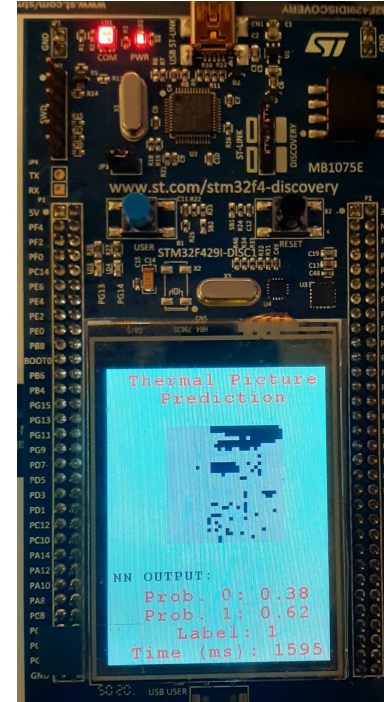


[https://community.arm.com/arm-community-blogs/b/architectures-and-processors-blog/posts/new-neural-network-kernels-boost-efficiency-in-microcontrollers-by-5x?utm\\_source=keil\\_com](https://community.arm.com/arm-community-blogs/b/architectures-and-processors-blog/posts/new-neural-network-kernels-boost-efficiency-in-microcontrollers-by-5x?utm_source=keil_com)

# Squeezenet Implementation (5)

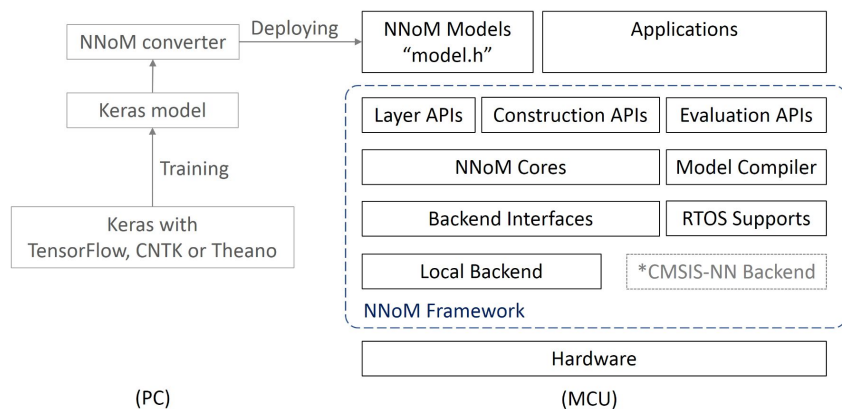
- STM TF lite C API
  - Experimental
  - No documentation
- Code glimpse / Demo

```
// Run inference
nn_status = tflm_c_invoke(hdl);
if(nn_status)
    _Exit(nn_status);
```



# Squeezenet Conclusion


- Maybe consider other toolchain:
  - C++ (with Mbed?)
  - NNoM <https://majianjia.github.io/nnom/>



# Gradient Boosted Trees (GBT) (1)

## General Approach

- Implementation of 2 Models:
  - a. Feature Generator - Convolutional Layers from CNN
  - b. Classifier - GBT (& SVC) Classifier
- Classification based on 512 Features from CNN



# Gradient Boosted Trees (GBT) (2)

## Implementation and Procedure


- Hyperparameter Tuning of CNN and Extraction of Convolutional Layers
- Hyperparameter Tuning of GBT
- Additional Approach: Support Vector Classifier (SVC) as Classifier on Same Features



# Classification Results


## GBT Classifier

Training Accuracy: 98.8%  
Test Accuracy: 70.0%



## Support Vector Classifier

Training Accuracy: 94.4%  
Test Accuracy: 70.0%



# GBT Implementation (1)

## Issues with CubeMX AI 7.0.0

- GBT Classifiers not supported → SVC
- Joined ONNX-Models not supported
- Multi-Output Models not supported
  - Classifier as ONNX-Model has two outputs:  
Labels & Probabilities
  - ONNX-Model Conversion → OUT OF SCOPE!

# GBT Implementation (2)

## Sizes of Networks on Board

Feature Generator →

Support Vector Classifier (SVC) →

Model manager			
Name	RAM	Flash	Complexity
network	10.28 KiB	151.88 KiB	78780 MACC
network_3	2.01 KiB	118.24 KiB	302080 MACC
<b>Total (2)</b>	<b>12.29 KiB (192.00</b>	<b>270.12 KiB (2.00</b>	<b>380860 MACC</b>

Approach would be RAM- and Flash-efficient!

# Conclusion (1)

## Data

- Labelling is highly subjective and ambiguous
- (Test) Dataset is very small compared to statistical practice (Here: 20 test samples vs. Good practice: 50 test samples and more)
- Class imbalance in the dataset

# Conclusion (2)

## Models and Implementation with CubeMX AI 7.0.0

- Some state-of-the-art models are not supported
- CNN has big size compared to other classifiers with comparable accuracy
- Implementation of non-standard models is workintense and out of scope for this project